

Assignment 3, Fall 2015

CS 4630, Defense Against the Dark Arts

x86 Assembly Language and Tools

Purpose

This assignment will refresh your knowledge of x86 assembly language (which you will have to analyze in all the virus detection assignments later in this semester), reinforce your knowledge of the C calling convention, and familiarize you with a couple of key tools used in the analysis of programs in their binary form (e.g., gdb). The assignment will also refresh your memory regarding writing, compiling, and running C code on Linux.

Due Date

This assignment is due on Wednesday, 09/30/2015, at 10:00am.

x86 Assembly Language Programming

This assignment will be completed on your Ubuntu installation.

To exercise your x86 skills, do the following:

1. Consider the following C code that calls function `dot_product()` (this C code is available on the Collab):

```
#include <stdio.h>

#define LOOP_COUNT 10000000
#define VSIZE 200

float v1[VSIZE], v2[VSIZE];
float dot_product(float *, float *, int length);

int main(int argc, char *argv[]) {
    float result;
    int i;

    /* initialize the two vectors */
    for (i = 0; i < VSIZE; i++) {
        v1[i] = (float) 2;
        v2[i] = (float) 3;
    }

    /* call dot_product some number of times for timing purposes */
    for (i = 0; i < LOOP_COUNT; i++)
        result = dot_product(v1, v2, VSIZE);

    printf ("result is %f\n", result);
    return 0;
}
```

2. For this assignment you will create an assembly language file containing a function called `dot_product()`. The function `dot_product()` computes the dot product of the two vectors (the first two parameters to function `dot_product()`) and returns the result. Because we are writing assembly language, we can hand-tune our code to make it as efficient as possible. Consequently, we should use the MMX/SSE extensions which provide “vector” instructions. For example, there is an instruction that will multiply two vectors consisting of four floating-point values.

CS4630 Assignment 3

3. To get started, you should write a C version of `dot_product()` and get the assembly code. The first step is to make sure our C code works properly. Here is the compilation :

```
gcc -m32 -march=corei7 -o a.out main.c dot_product.c
```

Remember, the option `-m32` tells gcc to generate code for the x86 (i.e., 32-bit code which we will always use in this class). The `-march=corei7` option tells the compiler we are using a machine that has the SSE3 extensions (unless you still have a Pentium, you will have these extensions).

Here is run of the program with the time command. You can check that the answer is correct.

```
time ./a.out
results is 1200.000000

real    0m6.585s
user    0m6.532s
sys     0m0.012s
```

Notice this code took about 6 seconds to run.

4. So the next step is to get the assembly code so we can modify it by hand. The appropriate command is:

```
gcc -m32 -c -S dot_product.c
```

You can use the code that the C compiler generated (in `dot_product.s`) as a starting point for your “optimized” code. That is, you should modify this assembly code to use the appropriate vector instructions. Here is a link to a short tutorial about the vector instructions:

http://neilkemp.us/src/sse_tutorial/sse_tutorial.html

There is much information on the web about MMX and SSE. Remember: Google is your friend (well most of the time). You can also read “Intel 64 and IA-32 Architectures Software Developer’s Manual” Volume 1 Chapter 10 and Volume 2 SSE instruction references, for more detailed explanation.

5. To run and compile your program using the assembly code file issue the following command:

```
gcc -m32 -march=corei7 -o a.out main.c dot_product.s
```

6. Here is a sample run:

```
time ./a.out
result is 1200.000000

0m1.269s
0m1.248s
```

CS4630 Assignment 3

0m0.008s

Notice how much faster the “vector” version ran.

Items to Submit

1. Submit your C program that you developed in step 3, `dot_product.c`, and your “optimized” assembly code `dot_product.s`.
2. It is mandatory that you use the file names specified to ease the task of grading 80+ submissions of this assignment. Throughout the semester, you will be given file names for assignments submitted using the class Collab Website.
3. I will give a five point bonus to anyone whose code runs faster than mine.